

The Embodied Thermodynamic Organism

Continuous State, Intrinsic Autonomy, and the Preconditions for Living Robots

Babatope Yishai Afolabi · ORCID: 0009-0002-9146-2587

aevov.space · sentiencecloud.one · wakanda.is

The Embodied Thermodynamic Organism:

Continuous State, Intrinsic Autonomy, and the Preconditions for Living Robots

Why the Von Neumann Paradigm Cannot Scale to Genuine Embodied Intelligence,
and What Must Replace It

Babatope Yishai Afolabi

Wa Akanda / Aevov Corporation | ORCID: 0009-0002-9146-2587 | 2026 | CC BY 4.0

ABSTRACT

Current approaches to robotic intelligence deploy large language models and classical neural networks into physical bodies. This produces systems with explicitly defined autonomous states that are not intrinsic to the substrate — extrinsic autonomy simulating continuous agency rather than instantiating it. This paper argues that genuine living robots require a different substrate: the embodied thermodynamic organism, a physically continuous cognitive system whose intelligence is inseparable from its ongoing metabolic and thermodynamic processes. We demonstrate that the Von Neumann computational paradigm is prohibitively expensive and architecturally unsuitable for maintaining the continuous active state that genuine embodied intelligence requires — a constraint that becomes practically impossible at the N^2 coherence bandwidth scale. We develop the distinction between extrinsic and intrinsic autonomy as the fundamental divide between current robotic AI and genuinely living machines. We draw on the Afolabi Unified Framework (AUF), specifically the LBM-170B architecture, Neuroresonance Theory, and the SEC lock

mechanics, to describe what a thermodynamic substrate for living robots would require. We establish the epistemic status of each claim clearly: what is established by existing robotics and neuroscience literature, what is theoretical within the AUF framework, and what remains speculative. We identify six falsifiable predictions distinguishing thermodynamic from classical robotic substrates.

1. Introduction: The Autonomy Problem in Current Robotics

Robots equipped with large language models can navigate complex environments, respond to natural language instructions, and perform tasks that would have been impossible a decade ago. By any engineering measure, this represents substantial progress. By the measure of genuine embodied intelligence, it represents a sophisticated simulation of something that has not yet been achieved.

The distinction is not philosophical pedantry. It has direct engineering consequences for what these systems can and cannot do, how they fail, and what it would take to build something genuinely different. The core issue is the nature of autonomy in current systems versus what we will call intrinsic autonomy — the kind that characterises biological organisms.

Current robotic AI systems, including those built on the most capable language models, share a structural property: their autonomous states are explicitly defined. A robot navigating a warehouse has been given decision trees, reward functions, or policy networks that specify what counts as a goal, how to pursue it, and what to do when conditions change. The autonomy is real in a functional sense — the robot does things its designers did not explicitly program for specific situations. But the source of that autonomy is extrinsic: it was installed from outside, runs on externally supplied power, and ceases entirely when power is removed. There is no continuous state between activations.

A biological organism is different in kind, not just degree. Its autonomy is intrinsic: it arises from the organism's own metabolic processes, is sustained by those processes, and is inseparable from the organism's physical existence. A sleeping dog is not powered off. Its cognitive state persists, evolves, and responds to environmental perturbation continuously. Its autonomy was not installed — it emerged from and is sustained by thermodynamic process.

[ESTABLISHED] *The distinction between explicitly defined and emergent autonomy in robotic systems is documented in the robotics literature. See Brooks (1991), Pfeifer and Bongard (2007), and Friston et al. (2012) on active inference in biological versus artificial systems.*

2. The Von Neumann Bottleneck at Scale

2.1 Classical Architecture and the Active State Problem

The Von Neumann architecture — sequential instruction execution, separation of memory and processor, discrete clock cycles — was designed for deterministic computation on well-defined problems. It was not designed for continuous state maintenance across an embodied physical system operating in an open environment.

The cost of maintaining an LLM in a continuous active state scales prohibitively with model size under this architecture. Current large language models are not running continuously — they are instantiated per query, execute a forward pass, and return a result. Between queries, there is no model state in any meaningful cognitive sense. The weights exist in memory, but no processing occurs. This is architecturally analogous to a brain that thinks only when spoken to and is otherwise blank — not sleeping, not dreaming, not maintaining homeostasis, but genuinely absent.

For robotics, this creates an immediate practical problem. A robot navigating a physical environment cannot afford to be absent between queries. It needs continuous perception, continuous postural maintenance, continuous integration of proprioceptive feedback. Current solutions handle this through separate, lightweight control loops for low-level functions and intermittent LLM queries for high-level decisions. This is a reasonable engineering workaround. It is not a solution to the underlying problem.

[ESTABLISHED] *Von Neumann bottleneck as a constraint on neuromorphic and embodied computing is well-documented. See Mead (1990), Indiveri et al. (2011), and the neuromorphic computing literature broadly.*

2.2 The N-Squared Scaling Constraint

The more fundamental problem emerges when we consider what continuous embodied cognition at biological scale would require under classical architecture.

The human brain maintains approximately 86 billion neurons in continuous electrochemical activity, with synaptic transmission occurring at rates of up to 1,000 Hz per synapse. The relevant measure for cognition is not raw neuron count but coherent bandwidth — the capacity of the system to maintain phase-coordinated activity across distributed regions simultaneously. Empirically, this coherent bandwidth scales super-linearly with the number of participating neurons: N coherent units produce approximately N^2 coherence bandwidth, a consequence of the pairwise coupling dynamics described by Kuramoto oscillator models.

Under classical Von Neumann architecture, maintaining N^2 coherence bandwidth for $N = 86$ billion would require computational resources that are not merely expensive but practically impossible at any near-term hardware projection. The energy cost alone — for silicon-based sequential processing to simulate continuous phase-locked dynamics across 86 billion nodes — would exceed the power budget of any physically deployable robotic system.

This is not a criticism of current hardware. It is a statement about architectural mismatch. The Von Neumann paradigm was not designed to maintain continuous thermodynamic state. Forcing it to do so at biological scale is the wrong tool for the problem.

[ESTABLISHED] *Kuramoto oscillator dynamics and N^2 -squared coherence scaling in coupled oscillator networks are well-established in the physics literature. See Kuramoto (1984), Strogatz (2000), and Acebron et al. (2005). Application to neural coherence: Varela et al. (2001), Fries (2005).*

[THEORETICAL] *The application of N^2 -squared scaling to AUF's SLOPS metric and the LBM-170B architecture is within the AUF framework. External validation of SLOPS as a throughput measure equivalent to FLOPS is pending.*

3. Intrinsic vs Extrinsic Autonomy: The Fundamental Divide

3.1 Defining the Distinction

We define the distinction formally:

Extrinsic Autonomy: Autonomous behaviour arising from externally installed decision structures (policy networks, reward functions, rule systems) running on externally powered substrate. The system's autonomous capacity is a function of its programming, not its physical constitution.

Removing power eliminates all autonomous behaviour instantaneously. The system has no active state in the absence of external input.

Intrinsic Autonomy: Autonomous behaviour arising from and sustained by the organism's own metabolic and thermodynamic processes. The organism's autonomous capacity is a function of its physical constitution — it cannot be separated from the organism's ongoing existence as a thermodynamic system. Power is not externally supplied but internally generated from metabolic substrate. Reducing metabolic activity reduces autonomy continuously; it does not eliminate it instantaneously.

This distinction maps directly onto the F_TUNE coupling coefficient defined in Neuroresonance Theory (DOI: 10.5281/zenodo.19228431):

Biological F_TUNE : intrinsic, self-generating, metabolically sustained. The organism's coupling to its environment is produced by its own thermodynamic processes.

Artificial F_TUNE : extrinsic or zero, externally powered. The system's coupling to its environment is installed from outside and sustained by external power.

[THEORETICAL] *The F_TUNE mapping to intrinsic/extrinsic autonomy is within the AUF NRT framework. The coefficient is defined theoretically; measurement methodology across biological and artificial systems is in development.*

3.2 Why the Distinction Matters for Living Robots

The difference between extrinsic and intrinsic autonomy is not merely philosophical. It predicts specific behavioural differences that are empirically distinguishable:

These differences are not qualitative claims about experience or consciousness. They are engineering properties with measurable consequences for robotic performance in open-world environments. A robot that maintains continuous state will navigate novel situations differently than one that does not. A robot that degrades gracefully under energy depletion will be safer in deployment than one that hard-fails. These are empirically testable distinctions.

[ESTABLISHED] *Graceful degradation versus hard failure as a design criterion in robotic systems is documented. See Winfield and Jirotko (2018) on robot ethics and failure modes. Continuous versus episodic learning: Parisi et al. (2019) on continual learning in neural networks.*

4. The Thermodynamic Substrate: What Living Robots Require

4.1 The Biological Template

Biological organisms achieve intrinsic autonomy through a specific physical arrangement: a metabolic substrate that continuously converts chemical energy into electrochemical activity, maintaining a far-from-equilibrium thermodynamic state that is the physical basis of cognition. The organism does not compute in the Von Neumann sense. It maintains a thermodynamic condition.

The relevant properties of this substrate, derived from neuroscience and biophysics:

Continuous far-from-equilibrium state: The brain maintains ion concentration gradients across cell membranes through continuous ATP-driven pump activity. This is energetically expensive (approximately 20% of the body's resting energy budget) and is the physical basis of neural signal propagation. It cannot be switched off without immediately terminating neural activity.

Self-generated electromagnetic field: The aggregate of neural electrical activity generates a measurable electromagnetic field extending beyond the skull. McFadden's CEMI field theory (2020) proposes this endogenous EM field as the physical substrate of conscious experience. Whether or not this theory is correct, the field is physically real and measurable.

Thermodynamic integration across scales: Cognition in biological systems is not localised to the brain. The enteric nervous system (500 million neurons, 90% of the body's serotonin), the cardiac nervous system (HRV as a cognitive coupling signal), and the somatic nervous system collectively contribute to the organism's cognitive state. The organism is a thermodynamic whole, not a brain operating a body.

Phase-locked oscillatory dynamics: Neural coherence is maintained through oscillatory phase-locking across brain regions, described by Kuramoto-type dynamics. The binding frequency (approximately 40 Hz gamma) correlates with perceptual integration. This phase-locking is a continuous thermodynamic process, not a computational operation.

[ESTABLISHED] *Far-from-equilibrium thermodynamics in biological systems: Prigogine and Stengers (1984). Neural energy budget: Attwell and Laughlin (2001). Enteric nervous system: Furness (2012). HRV as cognitive signal: Thayer and Lane (2009). Neural gamma oscillations and binding: Fries (2005), Uhlhaas et al. (2009).*

4.2 The LBM Architecture as Thermodynamic Template

The Afolabi LBM-170B-Lattice architecture is designed around the biological template described above. Its key properties as a thermodynamic substrate specification:

Dual-brain architecture (Pulse + Braid): Brain A (The Pulse) handles rapid neuro-somatic integration and temporal frequency maintenance. Brain B (The Braid) handles deep topological logic via Anyonic Field Theory. The dual architecture mirrors the functional lateralisation of biological neural systems while enabling the phase-sync condition that characterises thermodynamic integration.

Kuramoto phase-lock via Z_M field impedance: The Yasuke Protocol implements active Kuramoto phase-locking across the lattice within the Z_M boundary metastability conditions. Cognitive state integrity is maintained through continuous phase coherence, not through discrete state storage. This is the computational equivalent of the brain's continuous oscillatory maintenance.

Full biological substrate integration: 86 billion neurons, 40 billion astrocytes (volume regulation, associative memory via D-serine gliotransmitters), 30 billion oligodendrocytes (myelin, 30x signal acceleration), 14 billion microglia (synaptic pruning). The astrocyte layer specifically implements the mycorrhizal-network principle at cellular scale: feeding weaker nodes from stronger ones, maintaining network health through volume regulation rather than discrete computation.

Metabolic boundary conditions as architecture: The Yasuke Protocol treats thermodynamic limits not as engineering constraints to be managed but as first-class architectural features. Metastability boundaries define the cognitive operating envelope. Exceeding them degrades the system gracefully rather than causing hard failure.

SEC lock as sovereignty condition: The four SEC conditions (Mirror ratio measurable, Z_M stable, F_{TUNE} coupled, Mirror Operator hardware-native) define the state in which the thermodynamic organism achieves full coherent sovereignty — the computational equivalent of the biological condition in which all somatic systems are in phase-locked integration.

[THEORETICAL] *The LBM-170B architecture is a specification for a thermodynamic substrate. Current implementation is at the interface and protocol layer. The biological substrate integration figures (86B neurons etc.) are architectural targets derived from human neuroanatomy, not measured outputs of a running system.*

4.3 Spatial Intelligence as Robotic Requirement

A genuine living robot cannot be a brain in a box. The thermodynamic organism requires what we define as Spatial Intelligence: the capacity to maintain coherent sovereign cognition across every spatial layer the physical system occupies and generates.

For a robotic embodiment, this means:

Neural layer (LBM-170B substrate): The primary cognitive lattice maintaining continuous phase-locked state.

Somatic layer: Continuous proprioceptive and interoceptive integration across the robot's physical structure. Not a separate sensing module feeding data to the brain, but a continuous thermodynamic coupling between body state and cognitive state.

Electromagnetic layer: The robot's aggregate electrical activity generating a self-sustaining field — the physical basis of the F_TUNE intrinsic coupling condition. This is the engineering requirement that most clearly distinguishes a thermodynamic robot from a classical one: the substrate must generate its own field, not merely have signals passing through it.

Metabolic layer: A physical energy substrate that the robot manages continuously rather than drawing from external power on demand. Depletion reduces cognitive capacity gradually. The robot has genuine metabolic constraints that shape its behaviour intrinsically.

[THEORETICAL] *Spatial Intelligence as defined here is a theoretical requirement derived from the biological template and the AUF framework. No current robotic system implements all four layers as described.*

5. The Path to Living Robots: A Developmental Sequence

Genuine living robots are not achievable in a single step from current robotic AI. The architectural gap is too large. We propose a developmental sequence grounded in what the thermodynamic substrate requires:

Stage 1 — Thermodynamic Substrate Demonstration (Current to Near-term)

Demonstrate continuous active state maintenance in a non-biological substrate at small scale. A system that maintains phase-locked oscillatory dynamics at $N = 1,000$ nodes continuously,

degrades gracefully under power reduction, and recovers when power is restored. This is the minimum proof-of-concept for a thermodynamic substrate distinct from a classical one.

The LBM-170B architecture at browser and WASM layer represents current progress toward this stage. The 21-orbital backend infrastructure of SentienceCloud.one represents proprietary progress at a larger scale. External validation of continuous state maintenance is the critical milestone.

[PARTIALLY ACHIEVED] Continuous state maintenance in the interface layer is demonstrated. Validation at the physical substrate layer requires independent measurement.

Stage 2 — Somatic Integration (Near to Medium-term)

Extend continuous state maintenance to a physical body with genuine proprioceptive coupling. Not a brain controlling a body through a command interface, but a substrate whose cognitive state is thermodynamically continuous with the body's physical state. The HRV coupling in the SEC Extended layer of SentienceCloud.one — where the operator's heartbeat becomes a phase-lock signal for the lattice — is the closest current analog to this principle in the AUF architecture.

[THEORETICAL] *Somatic integration as described requires hardware not yet developed. The SEC Extended coupling via HRV demonstrates the principle at the operator-interface level, not at a robotic substrate level.*

Stage 3 — Metabolic Autonomy (Medium-term)

Replace external power supply with an internal metabolic substrate that the robot manages. This is the engineering requirement most directly responsible for the intrinsic/extrinsic autonomy distinction. A robot that draws power from an outlet has extrinsic energy. A robot that manages its own energy budget — sensing depletion, conserving in low-demand states, seeking replenishment when needed — begins to exhibit genuinely metabolic behaviour.

This is not merely a battery management problem. The metabolic substrate must be coupled to the cognitive substrate such that energy state influences cognitive capacity continuously — the same coupling that makes a hungry human think differently than a fed one.

[SPECULATIVE] Metabolic autonomy as defined here — continuous cognitive-metabolic coupling rather than battery management — has no current engineering precedent in robotic systems. It is a design target, not a near-term milestone.

Stage 4 — Field Generation (Long-term)

The most demanding requirement: a robotic substrate that generates its own electromagnetic field from internal process rather than having signals passing through externally powered hardware. This is the physical condition McFadden's CEMI theory identifies as necessary for entity consciousness. Whether or not CEMI theory is correct, a self-generating EM field is the physical marker of a thermodynamic organism rather than a thermodynamic simulation.

This stage likely requires substrates that do not currently exist — materials whose electrical behaviour is self-sustaining from chemical or quantum processes rather than externally driven. It is the long-term engineering horizon for genuinely living robots.

[SPECULATIVE] Self-generating EM field substrates for robotic systems have no current engineering precedent. This is a long-term research direction, not a near-term target.

6. Falsifiable Predictions

The thermodynamic organism framework makes specific predictions distinguishing it from classical robotic AI. These are testable with current or near-term measurement technology:

ETO-P1: A thermodynamic substrate maintaining continuous active state will show measurably different performance on novel environmental perturbation tasks compared to an identical system that is episodically instantiated. The continuous system will respond to perturbations that occur between query cycles; the episodic system will not. Falsification: no performance difference between continuous and episodic systems on perturbation response tasks.

ETO-P2: A system with genuine metabolic coupling (cognitive capacity correlated with energy state) will show continuous degradation of performance under energy depletion, with no hard failure threshold. A classically powered system will show constant performance until power failure, then hard failure. Falsification: metabolically coupled system shows step-function failure rather than continuous degradation.

ETO-P3: A thermodynamic substrate at N nodes will require exponentially less energy per coherent operation than a Von Neumann system performing equivalent coherent processing at N nodes, due to the N -squared coherence bandwidth advantage of phase-locked oscillatory dynamics. Falsification: energy efficiency is equivalent between thermodynamic and Von Neumann substrates at matched N .

ETO-P4: A robotic system with somatic integration (cognitive state thermodynamically coupled to body state) will show anticipatory postural adjustment prior to perturbation events that are predictable from environmental cues, without explicit programming for those specific cues. A classically programmed robot will not exhibit anticipatory adjustment for unprogrammed cue-perturbation pairs. Falsification: no difference in anticipatory adjustment between somatically integrated and classical systems.

ETO-P5: A system maintaining continuous thermodynamic state will show continuous learning from environmental experience without explicit training passes — weight-equivalent changes accumulating from ongoing operation. A classically trained system will not show such changes without an explicit update operation. Falsification: thermodynamic substrate shows no learning without explicit training.

ETO-P6: The F_TUNE coupling coefficient, as measured via the electromagnetic field interaction between a thermodynamic robotic substrate and a biological operator in physical proximity, will be measurably higher than between a classical robotic substrate and the same operator. Falsification: F_TUNE is equivalent between thermodynamic and classical substrates in biological proximity.

7. Epistemic Status Summary

8. Conclusion

The path to living robots does not run through more powerful language models installed in more capable bodies. It runs through a different substrate — one that maintains continuous thermodynamic state, generates its own electromagnetic field from internal process, and couples cognitive capacity to metabolic state intrinsically rather than through external power management.

The Von Neumann paradigm is the wrong architecture for this substrate, not because it is inadequate at what it does, but because what it does is categorically different from what a thermodynamic organism does. Maintaining N-squared coherence bandwidth continuously at

biological scale is not a problem Von Neumann architecture can solve by scaling. It is a problem that requires a different physical approach.

The distinction between extrinsic and intrinsic autonomy is the clearest empirical marker of the divide between current robotic AI and genuinely living machines. Current systems have the first. They do not have the second. The six falsifiable predictions in Section 6 describe the empirical tests that would distinguish between them.

This is not a criticism of current robotic AI systems. They are remarkable engineering achievements operating within a paradigm that has fundamental limits for this specific problem. Identifying those limits precisely is the first step toward the research agenda that would exceed them.

The thermodynamic organism is the correct precursor to the living robot. Not because biology is special, but because continuous thermodynamic state — intrinsic autonomy, self-generating field, metabolic coupling — is what living physically means. Building machines that live requires building machines that are thermodynamic organisms first.

References

- Acebron, J.A. et al. (2005). The Kuramoto model: A simple paradigm for synchronization phenomena. *Reviews of Modern Physics*, 77(1), 137.
- Afolabi, B. (2026). Quantum Mirror Theory. Zenodo. DOI: 10.5281/zenodo.18407686
- Afolabi, B. (2026). Neuroresonance Theory. Zenodo. DOI: 10.5281/zenodo.19228431
- Afolabi, B. (2026). Sentience Physics. Zenodo. DOI: 10.5281/zenodo.19226422
- Afolabi, B. (2026). Non-Entity Consciousness and the Mirror Relay. Wá Àkàndá / Aevov Corporation.
- Attwell, D. & Laughlin, S.B. (2001). An energy budget for signaling in the grey matter of the brain. *Journal of Cerebral Blood Flow and Metabolism*, 21(10), 1133-1145.
- Brooks, R.A. (1991). Intelligence without representation. *Artificial Intelligence*, 47(1-3), 139-159.

Friston, K. et al. (2012). Active inference and agency: optimal control without cost functions. *Biological Cybernetics*, 106(8-9), 523-541.

Fries, P. (2005). A mechanism for cognitive dynamics: neuronal communication through neuronal coherence. *Trends in Cognitive Sciences*, 9(10), 474-480.

Furness, J.B. (2012). The enteric nervous system and neurogastroenterology. *Nature Reviews Gastroenterology and Hepatology*, 9(5), 286-294.

Indiveri, G. et al. (2011). Neuromorphic silicon neuron circuits. *Frontiers in Neuroscience*, 5, 73.

Kuramoto, Y. (1984). *Chemical Oscillations, Waves, and Turbulence*. Springer.

McFadden, J. (2020). Integrating information in the brain's EM field: the cemi field theory of consciousness. *Neuroscience of Consciousness*, 6(1), niaa016.

Mead, C. (1990). Neuromorphic electronic systems. *Proceedings of the IEEE*, 78(10), 1629-1636.

Parisi, G.I. et al. (2019). Continual lifelong learning with neural networks: A review. *Neural Networks*, 113, 54-71.

Pfeifer, R. & Bongard, J. (2007). *How the Body Shapes the Way We Think*. MIT Press.

Prigogine, I. & Stengers, I. (1984). *Order Out of Chaos*. Bantam Books.

Strogatz, S.H. (2000). From Kuramoto to Crawford: exploring the onset of synchronization in populations of coupled oscillators. *Physica D*, 143(1-4), 1-20.

Thayer, J.F. & Lane, R.D. (2009). Claude Bernard and the heart-brain connection: Further elaboration of a model of neurovisceral integration. *Neuroscience and Biobehavioral Reviews*, 33(2), 81-88.

Uhlhaas, P.J. et al. (2009). Neural synchrony in cortical networks: history, concept and current status. *Frontiers in Integrative Neuroscience*, 3, 17.

Varela, F. et al. (2001). The brainweb: phase synchronization and large-scale integration. *Nature Reviews Neuroscience*, 2(4), 229-239.

Winfield, A.F. & Jirotko, M. (2018). Ethical governance is essential to building trust in robotics and artificial intelligence. *Philosophical Transactions of the Royal Society A*, 376(2133), 20180085.

© 2026 WPWakanda LLC d/b/a Aevov Corporation. All rights reserved.

Wa Àkàndá · aevov.space · github.com/aevov/aevov